

```
!pip install -q langchain_experimental langchain_core google-generativeai==0.3.1 google-ai-generativelanguage==0.4.0 langchain-google-gemini
!pip install -q "langchain[docarray]"
```

```
_____ 167.0/167.0 kB 4.4 MB/s eta 0:00:00
_____ 241.2/241.2 kB 12.6 MB/s eta 0:00:00
_____ 146.6/146.6 kB 13.6 MB/s eta 0:00:00
_____ 815.9/815.9 kB 18.0 MB/s eta 0:00:00
_____ 55.4/55.4 kB 5.7 MB/s eta 0:00:00
_____ 1.7/1.7 MB 28.2 MB/s eta 0:00:00
_____ 49.4/49.4 kB 4.1 MB/s eta 0:00:00
_____ 215.3/215.3 kB 5.2 MB/s eta 0:00:00
_____ 139.0/139.0 kB 7.6 MB/s eta 0:00:00
```

```
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
Building wheel for hnswlib (pyproject.toml) ... done
```

```
!pip -q install langchain huggingface_hub openai google-search-results tiktoken chromadb lark
```

```
_____ 226.7/226.7 kB 6.7 MB/s eta 0:00:00
Preparing metadata (setup.py) ... done
_____ 1.8/1.8 MB 28.2 MB/s eta 0:00:00
_____ 509.0/509.0 kB 17.3 MB/s eta 0:00:00
_____ 111.7/111.7 kB 9.7 MB/s eta 0:00:00
_____ 75.9/75.9 kB 8.0 MB/s eta 0:00:00
_____ 2.4/2.4 MB 40.7 MB/s eta 0:00:00
_____ 92.1/92.1 kB 10.4 MB/s eta 0:00:00
_____ 60.8/60.8 kB 6.9 MB/s eta 0:00:00
_____ 41.1/41.1 kB 4.4 MB/s eta 0:00:00
_____ 5.4/5.4 MB 56.2 MB/s eta 0:00:00
_____ 6.8/6.8 MB 79.1 MB/s eta 0:00:00
_____ 57.9/57.9 kB 6.2 MB/s eta 0:00:00
_____ 105.6/105.6 kB 11.5 MB/s eta 0:00:00
_____ 67.3/67.3 kB 8.0 MB/s eta 0:00:00
```

```
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
```

```
_____ 698.9/698.9 kB 42.1 MB/s eta 0:00:00
_____ 1.6/1.6 MB 59.6 MB/s eta 0:00:00
_____ 67.6/67.6 kB 7.6 MB/s eta 0:00:00
_____ 71.5/71.5 kB 8.6 MB/s eta 0:00:00
_____ 76.9/76.9 kB 8.4 MB/s eta 0:00:00
_____ 58.3/58.3 kB 5.7 MB/s eta 0:00:00
_____ 46.0/46.0 kB 5.2 MB/s eta 0:00:00
_____ 50.8/50.8 kB 5.6 MB/s eta 0:00:00
_____ 341.4/341.4 kB 30.2 MB/s eta 0:00:00
_____ 3.4/3.4 MB 87.0 MB/s eta 0:00:00
_____ 1.3/1.3 MB 61.2 MB/s eta 0:00:00
_____ 130.2/130.2 kB 13.2 MB/s eta 0:00:00
_____ 86.8/86.8 kB 10.5 MB/s eta 0:00:00
```

```
Building wheel for google-search-results (setup.py) ... done
Building wheel for pypika (pyproject.toml) ... done
```

```
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
lida 0.0.10 requires kaleido, which is not installed.
lida 0.0.10 requires python-multipart, which is not installed.
llmx 0.0.15a0 requires cohere, which is not installed.
```

```
!pip -q install sentence_transformers
!pip -q install -U FlagEmbedding
```

```
_____ 132.8/132.8 kB 4.1 MB/s eta 0:00:00
Preparing metadata (setup.py) ... done
_____ 536.6/536.6 kB 8.6 MB/s eta 0:00:00
_____ 280.0/280.0 kB 9.9 MB/s eta 0:00:00
_____ 38.3/38.3 MB 28.3 MB/s eta 0:00:00
_____ 116.3/116.3 kB 13.8 MB/s eta 0:00:00
_____ 134.8/134.8 kB 13.0 MB/s eta 0:00:00
```

```
Building wheel for FlagEmbedding (setup.py) ... done
```

```
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
ibis-framework 7.1.0 requires pyarrow<15,>=2, but you have pyarrow 15.0.0 which is incompatible.
```

```
import os
```

```
os.environ['GOOGLE_API_KEY']=""
```

```
!pip show langchain-google-genai
```



```
Name: langchain-google-genai
Version: 0.0.9
Summary: An integration package connecting Google's genai package and LangChain
Home-page: https://github.com/langchain-ai/langchain
Author:
```

```
Author-email:
License: MIT
Location: /usr/local/lib/python3.10/dist-packages
Requires: google-generativeai, langchain-core
Required-by:
```

```
!mkdir -p Data
!unzip -q /content/langchain_blog_posts.zip -d Data
```

```
import os
import google.generativeai as genai
```

```
genai.configure(api_key="")
```

```
from langchain.vectorstores import FAISS
```

```
from langchain.schema import Document
from langchain.vectorstores import Chroma
```

```
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain.storage import InMemoryByteStore
from langchain.document_loaders import TextLoader
```

```
from langchain_google_genai.embeddings import GoogleGenerativeAIEmbeddings
from langchain_google_genai.chat_models import ChatGoogleGenerativeAI
```

```
embeddings = GoogleGenerativeAIEmbeddings(model = "models/embedding-001")
```

```
loaders = [
    TextLoader('/content/Data/blog.langchain.dev_announcing-langsmith_.txt'),
    TextLoader('/content/Data/blog.langchain.dev_benchmarking-question-answering-over-csv-data_.txt'),
]
docs = []
for l in loaders:
    docs.extend(l.load())
```

```
len(docs)
```

```
2
```

```
text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=1000,
    chunk_overlap=200
)
```

```
texts = text_splitter.split_documents(docs)
```

```
# Helper function for printing docs
```

```
def pretty_print_docs(docs):
    print(f"\n{'-' * 100}\n".join([f"Document {i+1}:\n\n" + d.page_content for i, d in enumerate(docs)]))
```

```
!pip install -q faiss-cpu
```

```
----- 17.6/17.6 MB 45.2 MB/s eta 0:00:00
```

```
retriever = FAISS.from_documents(
    texts,
    embeddings,
).as_retriever()
```

```
docs = retriever.get_relevant_documents("What is LangSmith?")
```

```
pretty_print_docs(docs)
```

```
Document 1:
```

```
"Because we are building financial products, the bar for accuracy, personalization, and security is particularly high. LangSmith helps us ensure that we can't wait to bring these benefits to more teams. And we've got a long list of features on the roadmap like analytics, playground, and more."
-----
```

```
Document 2:
```

```
URL: https://blog.langchain.dev/announcing-langsmith/
```

```
Title: Announcing LangSmith, a unified platform for debugging, testing, evaluating, and monitoring your LLM applications
```

```
LangChain exists to make it as easy as possible to develop LLM-powered applications.
```

We started with an open-source Python package when the main blocker for building LLM-powered applications was getting a simple prototype

Document 3:

“Thanks to Langchain smith we were able to analyze our LLM calls, understand the performance of the different chain methods ( stuff

A unified platform

While each of these product areas provide unique value, often at a specific point in time in the development process, we believe a {

Document 4:

The blocker has now changed. While it’s easy to build a prototype of an application in ~5 lines of LangChain code, it’s still decept

Today, we’re introducing LangSmith, a platform to help developers close the gap between prototype and production. It’s designed for

LangSmith is now in closed beta. So if you’re looking for a robust, unified, system for debugging, testing, evaluating, and monitori

How did we get here?

## ✓ LLMChainExtractor + Contextual compression

```
from langchain_google_genai.chat_models import ChatGoogleGenerativeAI
from langchain.retrievers import ContextualCompressionRetriever
from langchain.retrievers.document_compressors import LLMChainExtractor
```

```
llm = ChatGoogleGenerativeAI(
    model="gemini-pro",
    temperature=0.8,
    convert_system_message_to_human=True
)
```

```
compressor = LLMChainExtractor.from_llm(llm)
```

```
compression_retriever = ContextualCompressionRetriever(
    base_compressor=compressor,
    base_retriever=retriever
)
```

```
compressor.llm_chain.prompt
```

```
PromptTemplate(input_variables=['context', 'question'], output_parser=NoOutputParser(), template='Given the following question and context, extract any part of the context *AS IS* that is relevant to answer the question. If none of the context is relevant return NO_OUTPUT. \n\nRemember, *DO NOT* edit the extracted parts of the context.\n\nQuestion: {question}\n\nContext:\n\n>>>\n{context}\n\n>>>\n\nExtracted relevant parts:')
```

```
compressed_docs = compression_retriever.get_relevant_documents("What is Langsmith?")
```

```
pretty_print_docs(compressed_docs)
```

```
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate
warnings.warn(
Document 1:
```

Today, we’re introducing LangSmith, a platform to help developers close the gap between prototype and production. It’s designed for

LangSmith is now in closed beta. So if you’re looking for a robust, unified, system for debugging, testing, evaluating, and monitori

◀ ▶

## ✓ LLMChainFilter

```
print("""Uses an LLM chain to select out the queries to show the final LLM - This could be shown to a model fine tuned to do this

"YES" we show it or "NO" we don't show it""")
```

Uses an LLM chain to select out the queries to show the final LLM - This could be shown to a model fine tuned to do this

"YES" we show it or "NO" we don't show it

```
from langchain.retrievers.document_compressors import LLMChainFilter
```

```
filter_ = LLMChainFilter.from_llm(llm)
```

```
filter_.llm_chain.prompt
```

```
PromptTemplate(input_variables=['context', 'question'], output_parser=BooleanOutputParser(), template="Given the following question and context, return YES if the context is relevant to the question and NO if it isn't.\n\n> Question: {question}\n\n> Context:\n>>>\n{context}\n\n>>>\n\n> Relevant (YES / NO):")
```

```
compression_retriever = ContextualCompressionRetriever(base_compressor=filter_, base_retriever=retriever)
```

```
compressed_docs = compression_retriever.get_relevant_documents("What is LangSmith")
```

```
pretty_print_docs(compressed_docs)
```

```
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate\n  warnings.warn(\n/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate\n  warnings.warn(\n/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate\n  warnings.warn(\n/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instead use predict_and_generate\n  warnings.warn(\nDocument 1:
```

“Because we are building financial products, the bar for accuracy, personalization, and security is particularly high. LangSmith helps us build better products faster. We can’t wait to bring these benefits to more teams. And we’ve got a long list of features on the roadmap like analytics, playground, and more.”

Document 2:

URL: <https://blog.langchain.dev/announcing-langsmith/>

Title: Announcing LangSmith, a unified platform for debugging, testing, evaluating, and monitoring your LLM applications

LangChain exists to make it as easy as possible to develop LLM-powered applications.

We started with an open-source Python package when the main blocker for building LLM-powered applications was getting a simple prototype. But now, the blocker has changed.

Document 3:

The blocker has now changed. While it’s easy to build a prototype of an application in ~5 lines of LangChain code, it’s still deceptively difficult to get to production.

Today, we’re introducing LangSmith, a platform to help developers close the gap between prototype and production. It’s designed for developers to debug, test, evaluate, and monitor their LLM applications.

LangSmith is now in closed beta. So if you’re looking for a robust, unified, system for debugging, testing, evaluating, and monitoring your LLM applications, we’ve got you covered.

How did we get here?

## ✓ EmbeddingsFilter

```
from langchain_google_genai.embeddings import GoogleGenerativeAIEmbeddings
```

```
from langchain_google_genai.chat_models import ChatGoogleGenerativeAI
```

```
embeddings = GoogleGenerativeAIEmbeddings(model = "models/embedding-001")
```

```
from langchain.retrievers.document_compressors import EmbeddingsFilter
```

```
embeddings_filter = EmbeddingsFilter(embeddings=embeddings, similarity_threshold=0.16)
```

```
compression_retriever = ContextualCompressionRetriever(base_compressor=embeddings_filter, base_retriever=retriever)
```

```
compressed_docs = compression_retriever.get_relevant_documents("What is LangSmith")
```

```
pretty_print_docs(compressed_docs)
```

Document 1:

“Because we are building financial products, the bar for accuracy, personalization, and security is particularly high. LangSmith helps us build better products faster. We can’t wait to bring these benefits to more teams. And we’ve got a long list of features on the roadmap like analytics, playground, and more.”

Document 2:

“Thanks to Langchain smith we were able to analyze our LLM calls, understand the performance of the different chain methods ( stuff we tried ) and choose the best one for our use case.”

A unified platform

While each of these product areas provide unique value, often at a specific point in time in the development process, we believe a {  
-----  
Document 3:  
  
URL: <https://blog.langchain.dev/announcing-langsmith/>  
Title: Announcing LangSmith, a unified platform for debugging, testing, evaluating, and monitoring your LLM applications  
  
LangChain exists to make it as easy as possible to develop LLM-powered applications.  
  
We started with an open-source Python package when the main blocker for building LLM-powered applications was getting a simple prot  
-----  
Document 4:  
  
The blocker has now changed. While it's easy to build a prototype of an application in ~5 lines of LangChain code, it's still decept  
Today, we're introducing LangSmith, a platform to help developers close the gap between prototype and production. It's designed for  
LangSmith is now in closed beta. So if you're looking for a robust, unified, system for debugging, testing, evaluating, and monitor:  
  
How did we get here?

## ✓ Pipelines

```
from langchain.document_transformers import EmbeddingsRedundantFilter
from langchain.retrievers.document_compressors import DocumentCompressorPipeline
from langchain.text_splitter import CharacterTextSplitter

splitter = CharacterTextSplitter(chunk_size=500, chunk_overlap=20, separator=". ")

redundant_filter = EmbeddingsRedundantFilter(embeddings=embeddings)
relevant_filter = EmbeddingsFilter(embeddings=embeddings, similarity_threshold=0.16)

pipeline_compressor = DocumentCompressorPipeline(
    transformers=[splitter, redundant_filter, relevant_filter]
)
```

```
compression_retriever = ContextualCompressionRetriever(base_compressor=pipeline_compressor,
    base_retriever=retriever)

compressed_docs = compression_retriever.get_relevant_documents("What is LangSmith")
pretty_print_docs(compressed_docs)
```

Document 1:  
  
URL: <https://blog.langchain.dev/announcing-langsmith/>  
Title: Announcing LangSmith, a unified platform for debugging, testing, evaluating, and monitoring your LLM applications  
  
LangChain exists to make it as easy as possible to develop LLM-powered applications.  
  
We started with an open-source Python package when the main blocker for building LLM-powered applications was getting a simple prot  
-----  
Document 2:  
  
It's designed for building and iterating on products that can harness the power-and wrangle the complexity-of LLMs.  
  
LangSmith is now in closed beta. So if you're looking for a robust, unified, system for debugging, testing, evaluating, and monitor:  
  
How did we get here?  
-----  
Document 3:  
  
“Thanks to Langchain smith we were able to analyze our LLM calls, understand the performance of the different chain methods ( stuff  
-----  
Document 4:  
  
We are consistently using it to improve our prompt engineering and look forward to the new features,” said Stan Girard, Head of Gen/  
  
A unified platform  
  
While each of these product areas provide unique value, often at a specific point in time in the development process, we believe a {  
-----  
Document 5:

We see teams with all kinds of Rube Goldberg-machine-like processes for managing their LLM applications, and we want to make that a

Document 6:

“Because we are building financial products, the bar for accuracy, personalization, and security is particularly high. LangSmith he

We can’t wait to bring these benefits to more teams. And we’ve got a long list of features on the roadmap like analytics, playground

Document 7:

The blocker has now changed. While it’s easy to build a prototype of an application in ~5 lines of LangChain code, it’s still decept

Today, we’re introducing LangSmith, a platform to help developers close the gap between prototype and production

Document 8:

We remember seeing Nat Friedman tweet in late 2022 that there was “not enough tinkering happening.” The LangChain open-source packag

```
pipeline_compressor = DocumentCompressorPipeline(
    transformers=[splitter, compressor, redundant_filter, relevant_filter]
)
```

```
compression_retriever = ContextualCompressionRetriever(base_compressor=pipeline_compressor,
    base_retriever=retriever)
```

```
compressed_docs = compression_retriever.get_relevant_documents("What is LangSmith")
pretty_print_docs(compressed_docs)
```

```
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instea
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instea
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instea
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instea
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instea
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instea
warnings.warn(
/usr/local/lib/python3.10/dist-packages/langchain/chains/llm.py:316: UserWarning: The predict_and_parse method is deprecated, instea
warnings.warn(
Document 1:
```

LangSmith is now in closed beta. So if you’re looking for a robust, unified, system for debugging, testing, evaluating, and monitori

Document 2:

A unified platform

While each of these product areas provide unique value, often at a specific point in time in the development process, we believe a g

Document 3:

Today, we’re introducing LangSmith, a platform to help developers close the gap between prototype and production

Document 4:

“Thanks to Langchain smith we were able to analyze our LLM calls, understand the performance of the different chain methods ( stuff

◀ ▶

Start coding or [generate](#) with AI.

